

Assignment 1: Imitation Learning

huseyinabanox@gmail.com

December 2022

1 Behavioral Cloning

1.1 Answer

Following files are modified:

- rl_trainer.py
- bc_agent.py
- MLP_policy.py
- replay_buffer.py
- utils.py
- pytorch_utils.py

1.2 Answer


eval_batch_size 5000
ep_len 1000
num_agent_train_steps_per_iter 1000
n_layers 2
size 64
learning_rate 0.0005

| | Hopper | Ant | HalfCheetah | Walker2d |
|---------------------|--------|-------|-------------|----------|
| Eval_AverageReturn | 961 | 4604 | 4009 | 441 |
| Eval_StdReturn | 285 | 91 | 67 | 325 |
| Train_AverageReturn | 3773 | 4714 | 4206 | 5567 |
| Train_StdReturn | 2 | 12 | 83 | 9 |
| Log_Loss | -351 | -1518 | -804 | -507 |

1.3 Answer

Walker2d is performing worst. Mean and std results indicate learning is not happening. First thing to play with is the learning rate. After finding a reasonable LR, then training steps are adjusted.

Behaviour Cloning - Walker2d



Network with 2 Hidden Layers of size 64, 10K training steps per iteration, LR=0.001

2 DAgger


2.1 Answer

DAgger implementation is similar.

2.2 Answer


Ant and Walker2D environments are run with DAgger. Training steps are set to 2K. Learning rate is set to 0.001.

DAgger - Walker2d



Network with 2 Hidden Layers of size 64, 2K training steps per iteration, LR=0.001

DAgger - Ant



Network with 2 Hidden Layers of size 64, 2K training steps per iteration, LR=0.001